

Software System Components Notes

School of Computer Science Students

October 5, 2015

Preface

This is a shared collection of notes for Software System Components. Please visit <https://www.github.com/UoB-CS-Students/Year-2-SSC> to find out more, and to see other modules. You can contribute to this document by:

- Editing the \LaTeX document you wish to contribute to, then submit a pull request to <https://www.github.com/UoB-CS-Students/Year-2-SSC>.
- Creating a new chapter, placing the \LaTeX file in the `tex` folder, and adding a line to `Notes.tex` such as `\input{./tex/MY_CHAPTER.tex}`

Here are some points to follow:

- For the purposes of version control, please try to put each sentence on a new line. (\LaTeX treats a single new line as a space, so inserting these extra spaces won't affect the display of your document).
- Place any package imports in `Notes.sty`.
- If you wish to contribute, try to make fairly small changes, and then submit a pull request.
- Use hyphens instead of spaces in your file names, e.g. `My-File.tex` instead of `My File.tex`
- Follow the current naming convention for files/chapters. For example, if the current file names are `1-Alpha`, `2-Beta`, then you should name your file `n-FILENAME`.

Contents

1	Introduction to Databases	1
1.1	How can we store (and communicate) data?	1
2	DBMS	2
2.1	What is a DBMS?	2
2.2	Relational databases	2
2.3	Database Models	3
2.4	What makes a good database?	3
2.5	ACID	3
2.6	ANSI-SPARC Architecture	4
2.7	Database applications	4
2.8	Data definition language (DDL)	4
2.9	Data manipulation language	4
2.10	Relational definitions	5
2.11	Keys	5
3	DML	6
3.1	Drop and alter	6
3.1.1	Drop	6
3.1.2	Alter	6
3.2	Types	6
4	Constraints	8
4.1	Keys	8
4.1.1	Primary Key	8
4.1.2	Foreign Key	9
4.2	Deletion	10
4.3	Domain Constraints	11
4.4	Constraints: Summary	11

Chapter 1

Introduction to Databases

1.1 How can we store (and communicate) data?

Store	Share
<ul style="list-style-type: none">• In memory (persistently?)• Java objects• Flat/structured data files:<ul style="list-style-type: none">– XML– Interchange formats	<ul style="list-style-type: none">• Shared memory• Share files locally• Share files remotely• Document repository• Data server—fetch data as required.

Chapter 2

DBMS

2.1 What is a DBMS?

A DBMS consists of a core database which provides:

1. Representation of data
2. Retrieval and maintenance of data

It also has a **collection of tools**:

- Design and build and maintenance
- E.g. standard applications

And it provides:

- An abstraction from implementation
- Efficient implementation
- Integrity, fault tolerance, security, . . .
- Standard (and bespoke) interfaces, e.g. SQL for querying a database

2.2 Relational databases

Relational databases are the pre-eminent choice for *general purpose* database systems, especially for commercial/enterprise systems since consistency, reliability etc. are critical. Here are some points about relational databases:

- Main model for database systems

- Data in form of sets and relations
- Data connected using basic set theory (selecting, combining, etc.)
- Normally viewed as **tables**
- Queries specify result, but now how it is computed, i.e. it is *declarative*

2.3 Database Models

We don't always need all the benefits of a relational database; we may not even need to write to or update our database. But we may need performance, scalability, and flexibility.

2.4 What makes a good database?

- It should be an *organised* collection of data
 - For a purpose
 - To facilitate some set of activities
 - Organised so that it can be accessed and maintained efficiently.
- Data independence from internal or physical representation.
- Minimise redundancy (only store data once)
- Maximise consistency (one underlying representation)
- Enable integration and sharing
- Facilitate change
- Logical organisation-

2.5 ACID

Atomicity a transaction happens as a whole, or not at all.

Consistency a database is always in a consistent state (according to the rules defined)

Isolation the effect of two operations happening in parallel is the same as if they had happened sequentially

Durability once something is stored it won't disappear

It should do all this even if the plug is unpulled.

2.6 ANSI-SPARC Architecture

ANSI-SPARC is an abstract design standard for a DBMS. It has three different levels:

External level (user views) a user's view of the database describes a part of the database that is relevant to a particular user.

Conceptual level is a way of describing what data is stored within the whole database and how the data is interrelated.

Internal level involves how the database is physically represented on the computer system. It describes how the data is actually stored in the database on the computer hardware.

2.7 Database applications

Users are usually shielded from the underlying database by application programs and web interfaces; this is for convenience and security.

2.8 Data definition language (DDL)

A data definition language is used to create, modify, and delete parts of the definition of the database.

Here is a sample of SQL code that would create a table called **student** with the fields **sid**, **dob**, **login**, and **course**.

```
CREATE TABLE Student (  
    sid      INTEGER  
    dob      CHAR(10)  
    login    CHAR(20)  
    course   CHAR(10)  
)
```

2.9 Data manipulation language

A data manipulation language is used to manipulate the data.

Here is a sample of SQL code that would select the **sid** and **login** columns from the **Student** table where the course name is **Se**:

```

SELECT  sid , login
FROM    Student
WHERE    course='Se '

```

We can impose constraints on the data. We can define these when we create the table, or we can add them later. Here is an example of imposing restraints:

```

CREATE TABLE Student (
    sid          INTEGER NOT NULL UNIQUE,
    dob          CHAR(10) ,
    login        CHAR(20) UNIQUE,
    course       CHAR(10)
)

```

2.10 Relational definitions

Domain an arbitrary (non-empty) set of atomic values.

Attribute name a symbol with an associated domain, $\text{dom}(A)$.

Relational schema a finite set of attribute names.

Tuple, t , of a relational schema R a mapping from attributes A of a relational schema R to the union of their domains $\text{dom}(A) : t(A) \in \text{dom}(A)$

Relation, r , of a relational schema R a finite set of tuples of relational schema R

Degree (Arity) of a relational schema the number of attributes

Cardinality of a relation the number of tuples

2.11 Keys

Superkey a set of attributes that can *always* be used to differentiate one tuple from another (within a relation).

Key a minimal superkey.

Concatenated key a key with more than one attribute.

Candidate key any key.

Primary key one of the candidate keys.

Foreign key an attribute of the relation which is the key for another relation.

Chapter 3

DML

3.1 Drop and alter

3.1.1 Drop

DROP deletes the table.

```
DROP TABLE Student
```

3.1.2 Alter

ALTER modifies the table definition.

```
ALTER TABLE Student  
  ADD COLUMN year_of_study INTEGER
```

3.2 Types

BOOLEAN TRUE, FALSE, or NULL

CHAR(size) or **CHARACTER(size)**

Strings Several versions

INTEGER or **INT** Several variations

FLOAT or **DOUBLE**

REAL, **NUMERIC**, or **DECIMAL**

DATE

TIME a time-of-day value

Chapter 4

Constraints

Constraints place restrictions on the values (or set of values) that can be inserted into the database. They are hard constraints, meaning the DBMS will check and enforce them. They are metadata used to make explicit domain constraints, e.g. an age must be non-negative, and they maintain the integrity of the database, not just locally to the table, e.g. only record marks for students who are registered.

Examples of constraints are domain constraints,
e.g. `INTEGER`, `NOT NULL`, `UNIQUE`.

4.1 Keys

4.1.1 Primary Key

By tagging an attribute as a primary key, it enforces the constraints `UNIQUE` and `NOT NULL`. It also signals to the DBMS that this is *probably* an attribute that can be used to uniquely identify tuples.

A primary key constraint can be seen in the following example:

```
CREATE TABLE Student (  
    sid INTEGER,  
    ...  
    PRIMARY KEY (sid)  
)
```

And a relation with a primary key can be denoted as `Student(sid, dob, login, course)`.

Primary Key Example

```
CREATE TABLE Student (  
    sid INTEGER,  
    dob CHAR(10),  
    login CHAR(20),  
    course CHAR(10),  
    PRIMARY KEY (sid)  
)
```

which is the same as

```
CREATE TABLE Student (  
    sid INTEGER,  
    dob CHAR(10),  
    login CHAR(20),  
    course CHAR(10),  
    CONSTRAINT StudentsKey PRIMARY KEY (sid)  
)
```

by naming the constraint, we are able to manipulate it later (we don't want rebuild the whole database!)

4.1.2 Foreign Key

A foreign key defines a link to another table. It usually specifies the primary key in that other table. Any row inserted into *this* table must satisfy the constraint that the foreign key in *this* table must be matched by a key in the other table.

Foreign Key Example

```
CREATE TABLE Student (  
    sid INTEGER,  
    dob CHAR(10),  
    login CHAR(20),  
    course CHAR(10),  
    PRIMARY KEY (sid)  
)  
  
CREATE TABLE Marks (  
    sid INTEGER,  
    cid INTEGER,  
    mark INTEGER,  
    FOREIGN KEY (sid) REFERENCES Student (sid),  
    FOREIGN KEY (cid) REFERENCES Course (cid)  
)
```

This isn't just checked on insertion; it's guaranteed to *always* be true.

4.2 Deletion

What happens if we want to delete a record? We can't have some fields referencing non-existent values in other tables. There are several operations to deal with deletion:

ON DELETE RESTRICT checks whether or not the deletion will put the system in an inconsistent state *before* the deletion is performed.

ON DELETE NO ACTION checks whether or not the deletion will put the system in an inconsistent state *after* the deletion is performed.

ON DELETE CASCADE propagates the deletion in one table throughout the other tables.

ON DELETE SET DEFAULT/NULL instead of deleting the record, it gives it a default value, e.g. 999999.

So the integrity of the database is maintained by forbidding the deletion (**RESTRICT**, **NO ACTION**), or by deleting rows that reference the deleted key (**CASCADE**).

Example of Deletion Control

```
CREATE TABLE Marks (
    sid INTEGER,
    cid INTEGER,
    mark INTEGER,

    FOREIGN KEY (sid) REFERENCES Student (sid),
        ON DELETE CASCADE,
        ON UPDATE CASCADE,
    FOREIGN KEY (cid) REFERENCES Course (cid)
)
```

The code indicates that

- changes to `sid` in the Students table propagate through to the Marks table.
- changes to `cid` in the Courses table are forbidden (the default action is **NO ACTION**).

4.3 Domain Constraints

We can enforce additional constraints:

- Local constraints
- Constraints over more than one table

For example:

```
CREATE TABLE Marks (  
    sid INTEGER,  
    cid INTEGER,  
    mark INTEGER,  
  
    FOREIGN KEY (sid) REFERENCES Student (sid),  
    ON DELETE CASCADE,  
    ON UPDATE CASCADE,  
    FOREIGN KEY (cid) REFERENCES Course (cid),  
    CHECK (mark >= 0 and mark <= 100)  
)
```

4.4 Constraints: Summary

- Constraints express metadata about our domain—they are declarative, so no need to write code.
- They are guaranteed to be enforced (not just checks on insertion).
- When designing, capture all *hard* domain constraints and express them in the database.
- Do not include constraints which are desirable but not inviolable, e.g. someone might not have a surname, so don't add the constraint **NOT NULL** to the surname field.